

STAT-S631

Assignment 10

John Koo

```
dp <- loadNamespace('dplyr')
import::from(magrittr, `%>%`, `%<>%`)
import::from(car, Anova, ncvTest)
library(ggplot2)
theme_set(theme_bw())
```

Problem 1

```
# get the data
robey.df <- read.table('~/.dev/stats-hw/stat-s631/Robey.txt') %>%
  dp$mutate(country = rownames(.))

# full model with region first
model.1 <- lm(tfr ~ region * contraceptors, data = robey.df)
summary(model.1)
```

Call:

```
lm(formula = tfr ~ region * contraceptors, data = robey.df)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.54546	-0.26527	-0.04661	0.34689	1.30579

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.832351	0.194090	35.202	< 2e-16 ***
regionAsia	-0.322375	0.563627	-0.572	0.570
regionLatin.Amer	-0.237356	0.520948	-0.456	0.651
regionNear.East	0.631733	0.632999	0.998	0.324
contraceptors	-0.054099	0.007718	-7.009	1.41e-08 ***
regionAsia:contraceptors	-0.003795	0.012389	-0.306	0.761
regionLatin.Amer:contraceptors	0.003136	0.012044	0.260	0.796
regionNear.East:contraceptors	-0.013920	0.016141	-0.862	0.393

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5732 on 42 degrees of freedom

Multiple R-squared: 0.8667, Adjusted R-squared: 0.8445

F-statistic: 39.01 on 7 and 42 DF, p-value: < 2.2e-16

```
anova(model.1)
```

Analysis of Variance Table

Response: tfr

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	3	44.304	14.768	44.9534	3.576e-13 ***
contraceptors	1	45.045	45.045	137.1158	8.226e-15 ***
region:contraceptors	3	0.365	0.122	0.3706	0.7746
Residuals	42	13.798	0.329		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
Anova(model.1)
```

Anova Table (Type II tests)

Response: tfr

	Sum Sq	Df	F value	Pr(>F)
region	1.677	3	1.7018	0.1812
contraceptors	45.045	1	137.1158	8.226e-15 ***
region:contraceptors	0.365	3	0.3706	0.7746
Residuals	13.798	42		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
# full model with contraception rate first
```

```
model.2 <- lm(tfr ~ contraceptors * region, data = robey.df)
```

```
summary(model.2)
```

Call:

```
lm(formula = tfr ~ contraceptors * region, data = robey.df)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.54546	-0.26527	-0.04661	0.34689	1.30579

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.832351	0.194090	35.202	< 2e-16 ***
contraceptors	-0.054099	0.007718	-7.009	1.41e-08 ***
regionAsia	-0.322375	0.563627	-0.572	0.570
regionLatin.Amer	-0.237356	0.520948	-0.456	0.651
regionNear.East	0.631733	0.632999	0.998	0.324
contraceptors:regionAsia	-0.003795	0.012389	-0.306	0.761
contraceptors:regionLatin.Amer	0.003136	0.012044	0.260	0.796
contraceptors:regionNear.East	-0.013920	0.016141	-0.862	0.393

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5732 on 42 degrees of freedom

Multiple R-squared: 0.8667, Adjusted R-squared: 0.8445

F-statistic: 39.01 on 7 and 42 DF, p-value: < 2.2e-16

```
anova(model.2)
```

Analysis of Variance Table

Response: tfr

Df	Sum Sq	Mean Sq	F value	Pr(>F)
----	--------	---------	---------	--------

```

contraceptors      1 87.672  87.672 266.8706 <2e-16 ***
region            3  1.677    0.559   1.7018 0.1812
contraceptors:region 3  0.365    0.122   0.3706 0.7746
Residuals         42 13.798    0.329
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
Anova(model.2)
```

Anova Table (Type II tests)

Response: tfr

	Sum Sq	Df	F value	Pr(>F)
contraceptors	45.045	1	137.1158	8.226e-15 ***
region	1.677	3	1.7018	0.1812
contraceptors:region	0.365	3	0.3706	0.7746
Residuals	13.798	42		

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Part a

Here, for simplicity, we will just subscript in the order of the model. So for example, β_1 corresponds to `regionAsia`, β_4 corresponds to `contraceptors`, β_7 corresponds to `regionNear.East:contraceptors`, etc.

Type I anova (`anova(model.1)`)

This test is sequential. So the first line tests $H_0 : \beta_i = 0$ for $0 < i \leq 7$ and $H_A : \beta_1 \neq 0$ and $\beta_2 \neq 0$ and $\beta_3 \neq 0$ since these all fall under `region`. In other words, we want to know if `region` adds anything to an intercept-only model. Since p is small, we can conclude that it does.

The second line tests $H_0 : \beta_i = 0$ for $3 < i \leq 7$ and $H_A : \beta_i \neq 0$ for all of $3 < i \leq 4$. In other words, given a model that just uses `region` and `contraceptors` without the interaction, does `contraceptors` add significant predictive power to the model? Since p is small, we can conclude that a model with both (but no interaction term) is significantly different from a model that just uses `region`.

The third line tests $H_0 : \beta_i = 0$ for $4 < i \leq 7$ and $H_A : \beta_i \neq 0$ for all $4 < i \leq 7$. In other words, we want to know if the interaction terms add anything to a model without interaction terms. Since p is close to 1, we can say that it does not.

Type II anova (`Anova(model.1)`)

In the first line, we test if a model containing `region` and the interaction term is significantly different from the full model. Here we remove the interaction term per the marginality principle. So $H_0 : \text{tfr} \sim \text{contraceptors}$ and $H_A : \text{tfr} \sim \text{contraceptors} + \text{region} + \text{contraceptors:region}$. Since p is large, we can say that the alternative model does not significantly add to the null model.

In the second line, we switch `contraceptors` and `region`. Since p is small, we can say that `contraceptors + region:contraceptors` does add to a model with just `region`.

In the third line, we test the non-interaction “parallel” model vs. the full model with interactions. $H_0 : \text{tfr} \sim \text{region} + \text{contraceptors}$, and $H_A : \text{tfr} \sim \text{region} + \text{contraceptors} + \text{region:contraceptors}$. Since p is large, we can fail to reject the null hypothesis and say that there is no significant addition to the model by adding the interaction term.

Part b

Type II anova is not sequential. It looks at each term separately (other than the marginality principle). So the order in which the covariates are written in the model call does not matter. On the other hand, type I anova is sequential so the order makes an effect. In both models, the same set of covariates are used, so the type II anova does not change, but since the order in which they are listed changes, the type I anova results differ.

Part c

From the type II anova tests (doesn't matter which one), we might say that **region** does not make a difference in the model (i.e., **contraceptors** already explains all of the variance in **tfr** that **region** can). So a model we might consider is:

```
model.3 <- lm(tfr ~ contraceptors, data = robey.df)
summary(model.3)
```

Call:

```
lm(formula = tfr ~ contraceptors, data = robey.df)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.5493	-0.3013	0.0254	0.3957	1.2021

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.875085	0.156860	43.83	<2e-16 ***
contraceptors	-0.058416	0.003584	-16.30	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5745 on 48 degrees of freedom

Multiple R-squared: 0.847, Adjusted R-squared: 0.8438

F-statistic: 265.7 on 1 and 48 DF, p-value: < 2.2e-16

```
anova(model.3, model.1)
```

Analysis of Variance Table

Model 1: tfr ~ contraceptors

Model 2: tfr ~ region * contraceptors

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	48	15.840				
2	42	13.798	6	2.0425	1.0362	0.4158

So we can conclude that **region** does not add to a model with **contraceptors**. This is the same result we obtained in Homework Assignment 8, where we made this decision largely based on visualizations of the data.

Problem 2

Note that $[W]_{ij} = 0 \forall i \neq j$, and $[W]_{ii} > 0 \forall i$ such that $0 < i \leq n$. Furthermore, $[W^{1/2}]_{ij} = \sqrt{[W]_{ij}} \forall i, j \leq n$. Then since W and $W^{1/2}$ are diagonal matrices, they are symmetric, i.e., $W^T = W$ and $(W^{1/2})^T = W^{1/2}$.

Since $W^{1/2}$ is a diagonal matrix, $W^{1/2}W^{1/2}$ is also diagonal. Note that $[W^{1/2}W^{1/2}]_{ij} = \sum_k \sqrt{w_{ik}w_{kj}} = \sqrt{w_{ii}w_{jj}}$ where $w_{ij} = [W]_{ij}$ and the other terms in the sum are zero since $w_{ij} = 0$ when $i \neq j$. Then $[W^{1/2}W^{1/2}]_{ij} = 0$ when $i \neq j$ and $[W^{1/2}W^{1/2}]_{ii} = w_{ii}$. Then $W^{1/2}W^{1/2} = W$.

Since $Y|X \sim \mathcal{N}(X\beta, \sigma^2 W^{-1})$, $\hat{\beta} = (X^T W X)^{-1} X^T W Y$. But if we start with the claim $\hat{\beta} = (X^{*T} X^*)^{-1} X^{*T} Y^*$ where $X^* = W^{1/2} X$ and $Y^* = W^{1/2} Y$:

$$\begin{aligned} & (X^{*T} X^*)^{-1} X^{*T} Y^* \\ &= ((W^{1/2} X)^T W^{1/2} X)^{-1} (W^{1/2} X)^T W^{1/2} Y \\ &= (X^T W^{1/2} W^{1/2} X)^{-1} X^T W^{1/2} W^{1/2} Y \\ &= (X^T W X)^{-1} X^T W Y \end{aligned}$$

Which is just our WLS estimator for $\hat{\beta}$.


Problem 3

[From ALR 7.6]

Part 1

```
stopping.df <- alr4::stopping

ggplot(stopping.df) +
  geom_point(aes(x = Speed, y = Distance)) +
  stat_smooth(aes(x = Speed, y = Distance),
              se = FALSE, method = 'lm')
```



$E[\text{Distance} \mid \text{Speed}]$ appears to change as we move along x . From the plot, we can see that it curves along the OLS line, starting above it, dipping below, and then increasing above it again.

Based on this scatterplot and our intuition of the problem, I believe it makes sense to exclude the intercept term. If a car is not moving (i.e., $\text{Speed} = 0$), then it cannot have a stopping distance, so $\text{Distance} = 0$, which eliminates the intercept term.

Part 2

```
# const.var.mod <- lm(Distance ~ Speed + I(Speed ** 2), data = stopping.df)
# const.var.mod <- lm(Distance ~ I(Speed ** 2) - 1, data = stopping.df)
const.var.mod <- lm(Distance ~ Speed + I(Speed ** 2) - 1, data = stopping.df)
summary(const.var.mod)
```

Call:

```
lm(formula = Distance ~ Speed + I(Speed^2) - 1, data = stopping.df)
```

Residuals:

Min	1Q	Median	3Q	Max
-22.298	-5.223	-0.259	4.198	27.771

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
Speed	0.576599	0.200804	2.871	0.00564	**
I(Speed^2)	0.062145	0.006904	9.001	9.83e-13	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.852 on 60 degrees of freedom

Multiple R-squared: 0.9644, Adjusted R-squared: 0.9632

F-statistic: 813.5 on 2 and 60 DF, p-value: < 2.2e-16

```
# add the predictions to the data
```

```
stopping.df %<>%
```

```
  dp$mutate(distance.hat = predict(const.var.mod, stopping.df)) %>%
```


```
  dp$mutate(resid.ols = Distance - distance.hat)
```

```
ggplot(stopping.df) +
```

```
  geom_point(aes(x = distance.hat, y = resid.ols)) +
```

```
  geom_abline(slope = 0, colour = 'red') +
```

```
  labs(x = 'OLS predictions', y = 'OLS residuals')
```




```
ggplot(stopping.df) +
```


```
  geom_point(aes(x = Speed, y = resid.ols)) +
```

```
  geom_abline(slope = 0, colour = 'red') +
```

```
  labs(y = 'OLS residuals')
```



```
ggplot(stopping.df) +  
  geom_point(aes(x = Speed ** 2, y = resid.ols)) +  
  geom_abline(slope = 0, colour = 'red') +  
  labs(y = 'OLS residuals', x = expression(Speed2))
```

```
# part a
fv.test <- ncvTest(const.var.mod)
fv.test
```

```
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 23.20677    Df = 1    p = 1.454844e-06
```

```
# part b
speed.test <- ncvTest(const.var.mod, ~ Speed)
speed.test
```

```
Non-constant Variance Score Test
Variance formula: ~ Speed
Chisquare = 23.47833    Df = 1    p = 1.263289e-06
```

```
# part c
speed.speed2.test <- ncvTest(const.var.mod, ~ Speed + I(Speed ** 2))
speed.speed2.test
```

```
Non-constant Variance Score Test
Variance formula: ~ Speed + I(Speed^2)
Chisquare = 23.57714    Df = 2    p = 7.590831e-06
```

```
# difference between b and c?
1 - pchisq(speed.speed2.test$ChiSquare - speed.test$ChiSquare, 1)
```

```
[1] 0.7532587
```

We do not gain any significant information by including both `Speed` and `Speed2` compared to just `Speed`.

Part 3

```
# build the WLS model
speed.var.mod <- lm(Distance ~ Speed + I(Speed ** 2) - 1,
                    weights = Speed ** -1,
                    data = stopping.df)
summary(speed.var.mod)
```

Call:

```
lm(formula = Distance ~ Speed + I(Speed^2) - 1, data = stopping.df,
    weights = Speed^-1)
```

Weighted Residuals:

Min	1Q	Median	3Q	Max
-4.0703	-1.4275	-0.1057	1.3737	5.0866

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
Speed	0.614461	0.158903	3.867	0.000274 ***
I(Speed^2)	0.060784	0.006141	9.898	3.14e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.997 on 60 degrees of freedom

Multiple R-squared: 0.958, Adjusted R-squared: 0.9566

F-statistic: 683.6 on 2 and 60 DF, p-value: < 2.2e-16

```
summary(const.var.mod)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
Speed	0.57659878	0.200803820	2.871453	5.638438e-03
I(Speed^2)	0.06214515	0.006904165	9.001110	9.828141e-13

```
summary(speed.var.mod)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
Speed	0.61446093	0.158903019	3.866893	2.737223e-04
I(Speed^2)	0.06078404	0.006140893	9.898240	3.143928e-14

Note that now our estimate for $\hat{\beta} = (X^T W X)^{-1} X^T W Y$ and $\text{var}(\hat{\beta}|X) = \sigma^2 (X^T X)^{-1} X^T W^{-1} X (X^T X)^{-1}$.

Part 4

```
# weight matrix
# it's diagonal since we only have one covariate
W <- diag(stopping.df$Speed ** -1)

# model matrix
X <- model.matrix(~ Speed + I(Speed**2), data = stopping.df)

beta.hat <- solve(t(X) %*% X) %*% t(X) %*% stopping.df$Distance

# estimate standard error
```

```

var.beta.hat <- sigma(const.var.mod) ** 2 *
  solve(t(X) %*% X) %*%
  (t(X) %*% solve(W) %*% X) %*%
  solve(t(X) %*% X)

```