

STAT-S676

Assignment 2

John Koo

```
library(ggplot2)
import::from(magrittr, `%>%`, set_colnames)
theme_set(theme_bw())
```

See source code here: <https://github.com/johneverettkoo/stats-hw>

Problem 1

Most of this code is from Dr. Womack's file `lm_mcmc.cpp` (from Canvas).

```
cpp.file <- '~/dev/stats-hw/stat-s676/package/src/lm_mcmc_horseshoe.cpp'
writeLines(readLines(cpp.file))
```

```
// [[Rcpp::depends(RcppArmadillo)]]
// [[Rcpp::depends(RcppGSL)]]

#include<iostream>
#include<armadillo>
#include<RcppArmadillo.h>
#include<gsl/gsl_rng.h>
#include<gsl/gsl_randist.h>

void pseudo_inverse_and_rank
(
  const arma::mat &W,
  arma::mat& p_inv,
  arma::mat& sqrt_pinv_mat,
  unsigned int& rank
)
{
  unsigned int dims = W.n_cols;
  arma::vec EVAL(dims);
  arma::mat EVEC(dims,dims);
  double tol = static_cast<double>(dims)*EVAL.max()*arma::datum::eps;
  arma::eig_sym(EVAL,EVEC,W);
  arma::uvec inds1 = arma::find(EVAL>tol);
  arma::uvec inds0 = arma::find(EVAL<=tol);
  EVAL.elem(inds0).fill(0);
  arma::vec EE=arma::ones<arma::vec>(inds1.n_elem);
  EVAL.elem(inds1)=EE/EVAL.elem(inds1);
  rank = inds1.n_elem;
  sqrt_pinv_mat = EVEC*arma::diagmat(arma::sqrt(EVAL));
  p_inv = sqrt_pinv_mat*sqrt_pinv_mat.t();
}
```

```

arma::vec rand_norm(const gsl_rng * r,unsigned int d){
  arma::vec out(d);
  for(unsigned int i=0;i<d;++i){
    out(i)=gsl_ran_ugaussian_ratio_method(r);
  }
  return(out);
}

// http://gallery.rcpp.org/articles/simulate-multivariate-normal/
arma::mat mvnrmArma(int n, arma::vec mu, arma::mat sigma) {
  int ncols = sigma.n_cols;
  arma::mat Y = arma::randn(n, ncols);
  return arma::repmat(mu, 1, n).t() + Y * arma::chol(sigma);
}

[[Rcpp::export]]
Rcpp::List lm_mcmc_horseshoe
(
  const arma::vec &y,
  const arma::mat &X_covariates,
  const unsigned int& N_sims,
  const unsigned long int& seed
)
{
  unsigned int p = X_covariates.n_cols;
  unsigned int n = X_covariates.n_rows;
  double p_double = static_cast<double>(p);
  double n_double = static_cast<double>(n);
  double gamma_shape = (n_double + p_double + 2.0) * 0.5;
  double gamma_rate;

  arma::mat X(n,p+1);
  X.col(0) = arma::ones<arma::vec>(n);
  X.cols(1,p) = X_covariates;

  // precompute stuff
  arma::mat XtX = X.t()*X;
  arma::vec Xty = X.t()*y;
  double yty = arma::as_scalar(y.t()*y);
  arma::mat XtX_pinv;
  arma::mat XtX_pinv_sqrt;
  unsigned int rank;
  pseudo_inverse_and_rank(XtX,XtX_pinv,XtX_pinv_sqrt,rank);

  // init
  arma::vec beta = XtX_pinv*Xty;
  // arma::vec beta_mean = beta*n_double/(1+n_double);
  double sigma_sq = yty - arma::as_scalar(Xty.t()*beta);
  arma::vec tau = arma::ones<arma::vec>(p+1);
  arma::vec lambda = arma::ones<arma::vec>(p+1);
  double phi = 1.0;
  double eta = 1.0;

  // outputs

```

```

arma::vec sigma_sq_out(N_sims);
arma::mat beta_out(N_sims, p+1);
arma::mat tau_out(N_sims, p+1);

// set up rng
const gsl_rng_type * T;
gsl_rng * r;
gsl_rng_env_setup();
T = gsl_rng_mt19937;
r = gsl_rng_alloc (T);
gsl_rng_set(r,seed);

for(unsigned int i = 0; i < N_sims; ++i) {
  // draw beta
  beta = mvnrmArma(
    1,
    arma::vectorise(arma::inv(XtX + arma::diagmat(tau)) * Xty),
    sigma_sq * arma::inv(XtX + arma::diagmat(tau))
  ).t();
  beta_out.row(i) = beta.t();

  // draw sigma_sq
  gamma_rate = arma::as_scalar(
    0.5 * dot(y - X * beta, y - X * beta) +
    dot(tau % beta, beta)
  ) + 1.0;
  sigma_sq = 1.0 / R::rgamma(gamma_shape, 1.0 / gamma_rate);
  sigma_sq_out(i) = sigma_sq;

  // draw tau and lambda
  for (unsigned int j = 0; j < p+1; ++j) {
    tau(j) = arma::as_scalar(
      R::rexp(1.0) /
      arma::as_scalar(beta(j) * beta(j) / 2.0 / sigma_sq + lambda(j) / 2.0)
    );
    lambda(j) = R::rexp(1.0) / arma::as_scalar(tau(j) / 2.0 + phi / 2.0);
  }
  tau_out.row(i) = tau.t();

  // draw phi
  phi = R::rgamma((p_double + 1.0) / 2.0,
    1.0 / (arma::sum(lambda) / 2.0 + eta / 2.0));

  // draw eta
  eta = R::rexp(phi / 2.0 + 0.5);
}

gsl_rng_free(r);

return(Rcpp::List::create(Rcpp::Named("beta")=beta_out,
  Rcpp::Named("sigma_sq")=sigma_sq_out,
  Rcpp::Named("tau")=tau_out));
}

```

Next, let's compile and install, and then see if it works by comparing it to R's built-in `lm` function:

```
Rcpp::compileAttributes('~dev/stats-hw/stat-s676/package')
devtools::install('~dev/stats-hw/stat-s676/package')

library(Stat676pack)

y <- iris$Petal.Width
X <- as.matrix(dplyr::select(iris, Petal.Length, Sepal.Width, Sepal.Length))
summary(lm(y ~ X))
```

Call:

```
lm(formula = y ~ X)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.60959	-0.10134	-0.01089	0.09825	0.60685

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.24031	0.17837	-1.347	0.18
XPetal.Length	0.52408	0.02449	21.399	< 2e-16 ***
XSepal.Width	0.22283	0.04894	4.553	1.10e-05 ***
XSepal.Length	-0.20727	0.04751	-4.363	2.41e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.192 on 146 degrees of freedom

Multiple R-squared: 0.9379, Adjusted R-squared: 0.9366

F-statistic: 734.4 on 3 and 146 DF, p-value: < 2.2e-16

```
out <- lm_mcmc_horseshoe(y, X, 10000, 676)
summary(as.data.frame(out$beta[9001:1000, ]))
```

V1		V2		V3		V4	
Min.	:-0.939445	Min.	:0.4101	Min.	:-0.0348	Min.	:-0.40323
1st Qu.:	-0.245581	1st Qu.:	0.4959	1st Qu.:	0.1467	1st Qu.:	-0.23876
Median :	-0.099362	Median :	0.5164	Median :	0.1869	Median :	-0.20168
Mean :	-0.135756	Mean :	0.5161	Mean :	0.1866	Mean :	-0.20090
3rd Qu.:	-0.003894	3rd Qu.:	0.5365	3rd Qu.:	0.2274	3rd Qu.:	-0.16407
Max.	: 0.473285	Max.	:0.6289	Max.	: 0.4418	Max.	: 0.02277

```
sigma(lm(y ~ X))
```

```
[1] 0.1919671
```

```
summary(out$sigma[901:1000])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.03835	0.04856	0.05150	0.05278	0.05562	0.07167

Problem 2

```
# load the data
load('~/.dev/stats-hw/stat-s676/pyrimidine.RData')
X <- unname(as.matrix(dplyr::select(D, -y)))
y <- D$y

# build MCMC model
out <- lm_mcmc_horseshoe(y, X, 10000, 314159)
```


We can see if any of them are significantly not 0.

```
beta.df <- as.data.frame(out$beta) %>%
  set_colnames(c('intercept', colnames(D)[-ncol(D)]))
tau.df <- as.data.frame(out$tau) %>%
  set_colnames(c('intercept', colnames(D)[-ncol(D)]))
sigma_sq <- out$sigma_sq

beta.df %>%
  # to long df
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  # make sure the labels are plotted in order
  dplyr::mutate(covariate = factor(covariate,
                                   levels = c(paste0('x', seq(26, 1)),
                                   'intercept')) %>%

  ggplot() +
  geom_boxplot(aes(x = covariate, group = covariate, y = estimate)) +
  geom_hline(yintercept = 0, colour = 'red') +
  coord_flip() +
  labs(title = 'MCMC results')
```

MCMC results




```

beta.df[9001:10000, ] %>% # just the last 1000
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  dplyr::mutate(covariate = factor(covariate,
                                   levels = c(paste0('x', seq(26, 1)),
                                   'intercept')) %>%

  ggplot() +
  geom_boxplot(aes(x = covariate, group = covariate, y = estimate)) +
  geom_hline(yintercept = 0, colour = 'red') +
  coord_flip() +
  labs(title = 'MCMC results (last 1000 results)')

```

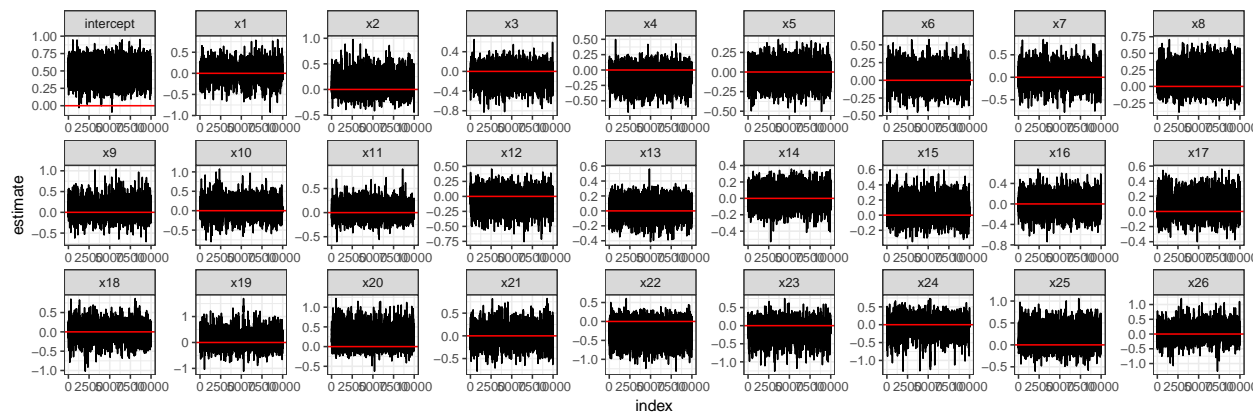
MCMC results (last 1000 results)




And we can see if each covariate found some steady-state-ish solution.

```
beta.df %>%
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  dplyr::mutate(covariate = factor(covariate,
                                   levels = c('intercept',
                                               paste0('x', seq(26)))),
               index = rep(seq(10000), 27)) %>%

  ggplot() +
  geom_line(aes(x = index, y = estimate)) +
  geom_hline(yintercept = 0, colour = 'red') +
  facet_wrap(~ covariate, scales = 'free', nrow = 3)
```



```
ggplot() +
  geom_line(aes(x = seq(length(sigma_sq)), y = sigma_sq)) +
  labs(x = 'index', y = expression(hat(sigma^2))) +
  coord_trans(y = 'log10')
```



It might be better to implement burn-in and thinning to account for the time it takes to converge as well as ACF:

```
## @title Horseshoe Gibbs sampler for linear models
## @description This is just a wrapper for the C++ code.
## @param y (numeric) An n-dimensional vector of responses
## @param X (numeric) An n by p dimensional matrix of inputs
## @param n.iter (numeric) The number of iterations to output
```

```

#' @param burn (numeric) Number of iterations for burn-in
#' @param thin (numeric) Thinning rate
#' @param seed (numeric) RNG seed (defaults to system time)
#' @return (list) The estimates for the covariates,  $\sigma^2$ , and tau
#' @export gibbs.horseshoe
gibbs.horseshoe <- function(y, X,
                           n.iter = 1e3, burn = 1e2, thin = 1e2,
                           seed = NULL) {
  # rng stuff
  if (is.null(seed)) seed <- as.numeric(Sys.time())

  # pass to C++ code
  out <- lm_mcmc_horseshoe(y, X, burn + n.iter * thin, seed)

  # take out burn-in
  beta. <- out$beta[-seq(burn), ]
  sigma_sq <- out$sigma_sq[-seq(burn)]
  tau <- out$tau[-seq(burn), ]

  # thin
  beta. <- beta.[seq(to = n.iter * thin, by = thin), ]
  sigma_sq <- sigma_sq[seq(to = n.iter * thin, by = thin)]
  tau <- tau[seq(to = n.iter * thin, by = thin), ]

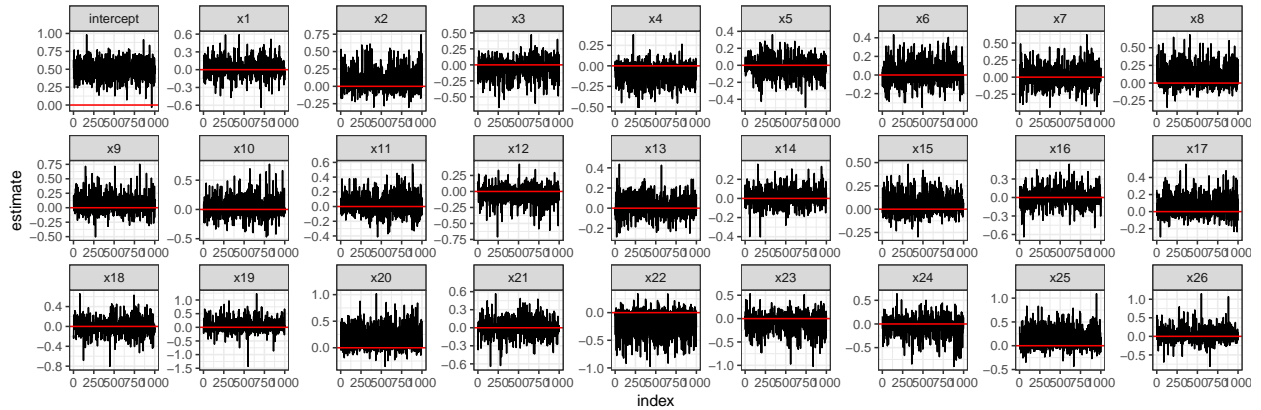
  return(list(beta = beta.,
              sigma_sq = sigma_sq,
              tau = tau))
}

out <- Stat676pack::gibbs.horseshoe(y, X, n.iter = 1e3)

beta.df <- as.data.frame(out$beta) %>%
  set_colnames(c('intercept', colnames(D)[-ncol(D)]))
tau.df <- as.data.frame(out$tau) %>%
  set_colnames(c('intercept', colnames(D)[-ncol(D)]))
sigma_sq <- out$sigma_sq

beta.df %>%
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  dplyr::mutate(covariate = factor(covariate,
                                  levels = c('intercept',
                                              paste0('x', seq(26)))),
               index = rep(seq(1e3), 27)) %>%
  ggplot() +
  geom_line(aes(x = index, y = estimate)) +
  geom_hline(yintercept = 0, colour = 'red') +
  facet_wrap(~ covariate, scales = 'free', nrow = 3)

```




We can also see how the τ s are distributed:

```
tau.df %>%
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  dplyr::mutate(covariate = factor(covariate,
                                   levels = c(paste0('x', seq(26, 1)),
                                               'intercept')) %>%

  ggplot() +
  geom_boxplot(aes(x = covariate, group = covariate, y = estimate)) +
  coord_flip() +
  labs(title = 'MCMC result') +
  scale_y_log10()
```

MCMC result



```

tau.df %>%
  tidyr::gather('covariate', 'estimate', 1:27) %>%
  dplyr::mutate(covariate = factor(covariate,
                                   levels = c('intercept',
                                               paste0('x', seq(26))))) %>%

  dplyr::group_by(covariate) %>%
  dplyr::summarise_all(median) %>%
  dplyr::ungroup() %>%
  ggplot() +
  geom_point(aes(x = covariate, y = estimate)) +
  geom_path(aes(x = covariate, y = estimate)) +
  labs(y = expression(tau))

```

